

✓ DATA COLLECTION AND PROCESSING

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
import requests
from bs4 import BeautifulSoup
from tqdm import tqdm
import re
import pandas as pd
```

```
def clean_text(text):
```

```
    text = re.sub(r'^(\Q\.[A-Z\s]+:)', '', text.strip())
```

```
    text = re.sub(r'^\w\s.?!\'\"', '', text)
```

```
    return text
```

```
def scrape_interview(interview_id):
```

```
    url = f"https://www.asapsports.com/show\_interview.php?id={interview\_id}"
```

```
    response = requests.get(url)
```

```
    if response.status_code == 200:
```

```
        soup = BeautifulSoup(response.content, 'html.parser')
```

```
        q = []
```

```
        a = []
```

```
        #print(soup.get_text())
```

```
        lines = soup.get_text().split('\n')
```

```
        question = False
```

```
        for line in lines:
```

```
            if 'Q.' in line:
```

```
                q.append(clean_text(line))
```

```
                question = True
```

```
            elif question == True and (':' in line):
```

```
                a.append(clean_text(line))
```

```
                question = False
```

```
        min_length = min(len(q), len(a))
```

```
        a = a[:min_length]
```

```
        q = q[:min_length]
```

```

        return q, a

    else:
        print(f"Failed ID: {interview_id}")
        return [], []

interview_ids = range(160001, 193174) # Example list of interview IDs

all_questions = []
all_answers = []
id_str = 0

for interview_id in tqdm(interview_ids):

    if interview_id % 10000 == 0:
        dict = {'questions': all_questions, 'answers': all_answers}

        df = pd.DataFrame(dict)

        id_str += 1
        #try:
        df.to_csv('./QAData' + str(id_str) + '.csv', index=False)
        #except:
        # df.to_csv('/content/drive/QAData' + str(id_str) + '.csv', index=False)

    questions, answers = scrape_interview(interview_id)
    all_questions.extend(questions)
    all_answers.extend(answers)

#for i in range(len(all_questions)):
#    print(f"Question {i + 1}: {all_questions[i]}")
#    print(f"Answer {i + 1}: {all_answers[i]}")
#    print()

dict = {'questions': all_questions, 'answers': all_answers}

df = pd.DataFrame(dict)

try:
    df.to_csv('/content/drive/My Drive/QAData.csv', index=False)
except:
    df.to_csv('/content/drive/QAData.csv', index=False)

46%|██████████| 15408/33173 [1:42:42<1:36:27, 3.07it/s]WARNING:bs4.dammit:Sc
100%|██████████| 33173/33173 [3:40:56<00:00, 2.50it/s]

#final_0_1 = pd.read_csv('./QAData0_1.csv')

```

```

#final_0_2 = pd.read_csv('./QAData0_2.csv')
#final = pd.concat([final_0_1, final_0_2], ignore_index=True)

df1 = pd.read_csv('./QAData1.csv')
df2 = pd.read_csv('./QAData2.csv')
df3 = pd.read_csv('./QAData3.csv')
df4 = pd.read_csv('./QAData4.csv')
df5 = pd.read_csv('./QAData5.csv')
df6 = pd.read_csv('./QAData6.csv')
df7 = pd.read_csv('./QAData7.csv')

final = pd.concat([df1, df2, df3, df4, df5, df6, df7])

all_questions = final['questions'].tolist()
all_answers = final['answers'].tolist()

print(len(all_questions))
print(len(all_answers))
for i in range(5):
    print("Q:", all_questions[i])
    print("A:", all_answers[i])
    print()
    344072
    344072
    Q: You're back here in this place where you've won before. Indiscernible w
    A: I love the golf course but the fans are what makes it awesome. To be abl

    Q: I caught a glimpse of you looking at the wall of achievements. I can't i
    A: Yeah it's an honor. A lot of names up there that I would've loved to pla

    Q: In terms of making your comeback now can you give us an update on how you
    A: I'm just taking it one day at a time really. It's been a struggle. I'm i

    Q: How tough is it to fight through that when you're trying to find the cons
    A: Yeah it's pretty tough. It's pretty tough to find. I don't know man. I

    Q: Does it give you a different perspective having to fight through it this
    A: Oh yeah. It gives me a lot bigger perspective especially my kids and stu

```

✓ Chat QA with GPT-2

```
!pip install transformers
```

```

Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: huggingface-hub<1.0,>=0.16.4 in /usr/local/lib
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: nvidia-ml>=5.1 in /usr/local/lib/python3.10/dist-

```

```

Requirement already satisfied: pyyaml<5.4,>=3.1 in /usr/local/lib/python3.10/dist-packages (5.3.1)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (2022.10.31)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (2.28.1)
Requirement already satisfied: tokenizers<0.19,>=0.14 in /usr/local/lib/python3.10/dist-packages (0.15.1)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.10/dist-packages (0.3.4)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (4.64.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (2023.9.2)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (4.5.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (2.0.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (2023.7.22)

```

```
from transformers import GPT2Tokenizer, GPT2LMHeadModel
```

```
import torch
```

```
device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
```

```
tokenizer.add_special_tokens({"pad_token": "<pad>",
                             "bos_token": "<startofstring>",
                             "eos_token": "<endofstring>"})
```

```
tokenizer.add_tokens(["<ans>:"])
```

```
model = GPT2LMHeadModel.from_pretrained("gpt2")
```

```
model.resize_token_embeddings(len(tokenizer))
```

```
model = model.to(device)
```

```

vocab.json: 100% 1.04M/1.04M [00:00<00:00, 7.08MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 4.42MB/s]
tokenizer.json: 1.36M/1.36M [00:00<00:00, 15.1MB/s]
100%
config.json: 100% 665/665 [00:00<00:00, 6.77kB/s]
model.safetensors: 548M/548M [00:06<00:00, 98.7MB/s]
100%

```

```
from torch.utils.data import Dataset
```

```
class QADData(Dataset):
```

```
    def __init__(self, tokenizer, q, a):
```

```
        self.X = []
```

```
        for i in range(len(q)):
```

```
            self.X.append("<startofstring> " + str(q[i]) + " <ans>: " + str(a[i]) + " <endofstring>")
```

```
    print(self.X[0])
```

```
    self.X_encoded = tokenizer(self.X, max_length=128, truncation=True, padding='max_length')
```

```

        self.X_encoded = tokenizer(self.X, max_length=128, truncation=True, padding='max_length')
        self.input_ids = self.X_encoded['input_ids']
        self.attention_mask = self.X_encoded['attention_mask']

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return (self.input_ids[idx], self.attention_mask[idx])

ds = QAData(tokenizer, all_questions, all_answers)
<startofstring> You're back here in this place where you've won before. Ind:

import numpy as np
np.save('gpt2_ds', ds)
data = np.load('gpt2_ds')

/usr/local/lib/python3.10/dist-packages/numpy/lib/npio.py:521: FutureWarning
    arr = np.asanyarray(arr)
/usr/local/lib/python3.10/dist-packages/numpy/lib/npio.py:521: VisibleDeprecationWarning
    arr = np.asanyarray(arr)

from torch.utils.data import DataLoader
QAData = DataLoader(ds, batch_size=16)

from tqdm import tqdm
import torch

def train(QAData, model, optim):

    epochs = 12

    for i in tqdm(range(epochs)):
        for X, a in QAData:
            X = X.to(device)
            a = a.to(device)
            optim.zero_grad()
            loss = model(X, attention_mask=a, labels=X).loss
            loss.backward()
            optim.step()
            torch.save(model.state_dict(), "model_state.pt")
            print(infer("how did you guys play today?"))

def infer(inp):
    inp = "<startofstring> "+inp+" <ans>: "
    inp = tokenizer(inp, return_tensors="pt")
    X = inp["input_ids"].to(device)
    a = inp["attention_mask"].to(device)
    output = model.generate(X, attention_mask=a )
    output = tokenizer.decode(output[0])

```

```

return output

from torch.optim import Adam
model.train()

optim = Adam(model.parameters(), lr=1e-3)

print("training .... ")
train(QADData, model, optim)

print("infer from model : ")
while True:
    inp = input()
    print(infer(inp))

    training ....
    0%|          | 0/12 [00:00<?, ?it/s]Setting `pad_token_id` to `eos_token_id`
    8%|█         | 1/12 [00:13<02:26, 13.33s/it]
    <startofstring> how did you guys play today? <ans>:
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    17%|█        | 2/12 [00:25<02:08, 12.83s/it]
    <startofstring> how did you guys play today? <ans>:
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    25%|██       | 3/12 [00:38<01:54, 12.70s/it]
    <startofstring> how did you guys play today? <ans>:          I mean we
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    33%|███      | 4/12 [00:51<01:41, 12.71s/it]
    <startofstring> how did you guys play today? <ans>:          I don't think
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    42%|████     | 5/12 [01:03<01:29, 12.73s/it]
    <startofstring> how did you guys play today? <ans>:          I played five time
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    50%|█████    | 6/12 [01:16<01:16, 12.73s/it]
    <startofstring> how did you guys play today? <ans>:          I mean they got a
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    58%|█████    | 7/12 [01:29<01:03, 12.76s/it]
    <startofstring> how did you guys play today? <ans>:          Yeah. I played in
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    67%|██████   | 8/12 [01:41<00:50, 12.68s/it]
    <startofstring> how did you guys play today? <ans>:          I mean it was a big
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    75%|██████   | 9/12 [01:54<00:37, 12.66s/it]
    <startofstring> how did you guys play today? <ans>:          I played five times :
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    83%|███████  | 10/12 [02:06<00:25, 12.58s/it]
    <startofstring> how did you guys play today? <ans>:          I played five years :
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    92%|███████  | 11/12 [02:19<00:12, 12.72s/it]
    <startofstring> how did you guys play today? <ans>:          I mean it's not a
    Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
    100%|███████ | 12/12 [02:32<00:00, 12.70s/it]
    <startofstring> how did you guys play today? <ans>:          I played five years

```

```
infer from model :
Coach what do you think of the team's performance?
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
<startofstring> Coach what do you think of the team's performance? <ans>:
Are you proud of the effort displayed today?
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
<startofstring> Are you proud of the effort displayed today? <ans>:      I'm
```

```
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-12-685f3796ca7b> in <cell line: 10>()
      9 print("infer from model : ")
     10 while True:
--> 11     inp = input()
     12     print(infer(inp))
```

1 frames

```
/usr/local/lib/python3.10/dist-packages/ipykernel/kernelbase.py in
```

✓ BERT For NSP

```
from transformers import BertTokenizer, BertForNextSentencePrediction
import torch
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertForNextSentencePrediction.from_pretrained('bert-base-uncased')
```

```
text = (all_questions[0])
text2 = (all_answers[0])
```

```
inputs = tokenizer(text, text2, return_tensors='pt')
inputs.keys()
```

```
tokenizer_config.json: 28.0/28.0 [00:00<00:00,
100% 658B/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 1.39MB/s]
tokenizer.json: 100% 466k/466k [00:00<00:00, 5.50MB/s]
config.json: 100% 570/570 [00:00<00:00, 29.8kB/s]
model.safetensors: 440M/440M [00:01<00:00,
```

```
inputs
```

```
{'input_ids': tensor([[ 101,  2074,  2115,  4301,  2006,  1996,  2457,
 2041,  2045,  1996,
                    6891,  1998,  2115,  4301,  2006,  2652,  2122,  2399,  1996,
 2345,
                    2261,  2399,  1999,  5869,  7136,  1012,   102,  2026,  4301,
```

```
sentence_a = []
sentence_b = []
```



```

label = []

for i in (range(len(all_questions))):
    # 50/50 whether is IsNextSentence or NotNextSentence
    if random.random() >= 0.2:
        # this is IsNextSentence
        sentence_a.append(all_questions[i])
        sentence_b.append(all_answers[i])
        label.append(0)
    else:
        index = random.randint(0, len(all_questions)-1)
        # this is NotNextSentence
        sentence_a.append(all_questions[i])
        sentence_b.append(all_answers[index])
        label.append(1)

for i in range(3):
    print(label[i])
    print(sentence_a[i] + '\n---')
    print(sentence_b[i] + '\n')

    0
    Just your thoughts on the court out there the venue and your thoughts on play:
    ---
    My thoughts on the court is it looks like a stage. That looks dope. This

    0
    BI you just mentioned how you guys aren't on national TV a lot. CJ reference
    ---
    Goofy. We come to work we all have fun. Off the court on the court we al

    0
    You mentioned that you guys like to have fun. Have you thought about how you
    ---
    No we ain't thought about it but I know it's going to be a good time.

inputs = tokenizer(sentence_a, sentence_b, return_tensors='pt', max_length=512, t

inputs.keys()
dict_keys(['input_ids', 'token_type_ids', 'attention_mask'])

inputs.token_type_ids[0]
tensor([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,
1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.

```

12/23/23, 10:33 PM

```

def __init__(self, encodings):
    self.encodings = encodings
def __getitem__(self, idx):
    return {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
def __len__(self):
    return len(self.encodings.input_ids)

```

```
dataset = QADataset(inputs)
```

```
loader = torch.utils.data.DataLoader(dataset, batch_size=8, shuffle=True)
```

```
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
# and move our model over to the selected device
model.to(device)
```

```

BertForNextSentencePrediction(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
          (intermediate): BertIntermediate(
            (dense): Linear(in_features=768, out_features=3072, bias=True)
            (intermediate_act_fn): GELUActivation()
          )
          (output): BertOutput(
            (dense): Linear(in_features=3072, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
    )
  )
)

```

```

    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (cls): BertOnlyNSPHead(
    (seq_relationship): Linear(in_features=768, out_features=2, bias=True)
  )
)

```

```
from transformers import AdamW
```

```
model.train()
```

```
optim = AdamW(model.parameters(), lr=5e-5)
```

```

/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411: FutureWarning:
warnings.warn(

```

```
import matplotlib.pyplot as plt
```

```
from tqdm import tqdm # for our progress bar
```

```
epochs = 5
```

```
y = []
```

```
for epoch in range(epochs):
```

```
    # setup loop with TQDM and dataloader
```

```
    loop = tqdm(loader, leave=True)
```

```
    for batch in loop:
```

```
        # initialize calculated gradients (from prev step)
```

```
        optim.zero_grad()
```

```
        # pull all tensor batches required for training
```

```
        input_ids = batch['input_ids'].to(device)
```

```
        attention_mask = batch['attention_mask'].to(device)
```

```
        token_type_ids = batch['token_type_ids'].to(device)
```

```
        labels = batch['labels'].to(device)
```

```
        # process
```

```
        outputs = model(input_ids, attention_mask=attention_mask,
```

```
                        token_type_ids=token_type_ids,
```

```
                        labels=labels)
```

```
        # extract loss
```

```
        loss = outputs.loss
```

```
        # calculate loss for every parameter that needs grad update
```

```
        loss.backward()
```

```
        # update parameters
```

```
        optim.step()
```

```
        # print relevant info to progress bar
```

```
        loop.set_description(f'Epoch {epoch}')
```

```
        loop.set_postfix(loss=loss.item())
```

```
        y.append(float(loss.item()))

```

```

        y.append(float(loss.item()))

plt.xlabel('Epochs')
plt.ylabel('Loss Per Epoch')
plt.title('Loss')

x = list(range(1, epochs + 1))

plt.plot(x, y)

plt.show()
0%|          | 0/130 [00:00<?, ?it/s]<ipython-input-10-834994efa95d>:5: Use
return {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
Epoch 0: 100%|          | 130/130 [00:32<00:00,  4.00it/s, loss=0.205]
Epoch 1: 100%|          | 130/130 [00:30<00:00,  4.22it/s, loss=0.00706]
Epoch 2: 100%|          | 130/130 [00:31<00:00,  4.08it/s, loss=0.0711]
Epoch 3: 100%|          | 130/130 [00:31<00:00,  4.18it/s, loss=0.0181]
Epoch 4: 100%|          | 130/130 [00:30<00:00,  4.20it/s, loss=0.00349]

# Regular BERT model for comparison
bert_model = BertForNextSentencePrediction.from_pretrained('bert-base-uncased')

bert_model = bert_model.to(device)

# Text input
input_text = "Do you learn something from it?"
next_sentence = "Yeah I think you learn"

# Tokenize the input
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
inputs = tokenizer(input_text, next_sentence, return_tensors='pt')

inputs = inputs.to(device)

# Get predictions from your fine-tuned model
with torch.no_grad():
    outputs = model(**inputs)
    prediction = "IsNext" if outputs.logits[0][0] > outputs.logits[0][1] else "No"

# Decode the input IDs to get the actual sentences
decoded_input = tokenizer.decode(inputs['input_ids'][0], skip_special_tokens=True)
decoded_next_sentence = tokenizer.decode(inputs['input_ids'][1], skip_special_tokens=True)

# Get predictions from the regular BERT model
with torch.no_grad():
    outputs_bert = bert_model(**inputs)
    prediction_bert = "IsNext" if outputs_bert.logits[0][0] > outputs_bert.logits[0][1] else "No"

# Decode the input IDs to get the actual sentences
decoded_input_bert = tokenizer.decode(inputs['input_ids'][0], skip_special_tokens=True)
decoded_next_sentence_bert = tokenizer.decode(inputs['input_ids'][1], skip_special_tokens=True)

```

```

decoded_next_sentence_bert = tokenizer.decode(inputs[1:input_ids[-1]], skip_special_tokens=True)

# Print the actual sentences and predictions
print(f"Fine-tuned Model Prediction: {prediction}")
print(f"Fine-tuned Model Input: {decoded_input}")
print(f"Fine-tuned Model Next Sentence: {decoded_next_sentence}")

print(f"\nRegular BERT Model Prediction: {prediction_bert}")
print(f"Regular BERT Model Input: {decoded_input_bert}")
print(f"Regular BERT Model Next Sentence: {decoded_next_sentence_bert}")

Fine-tuned Model Prediction: IsNext
Fine-tuned Model Input: do you learn something from it? yeah i think you learn
Fine-tuned Model Next Sentence: do you learn something from it? yeah i think you learn

Regular BERT Model Prediction: IsNext
Regular BERT Model Input: do you learn something from it? yeah i think you learn
Regular BERT Model Next Sentence: do you learn something from it? yeah i think you learn

```

✓ FAILED ATTEMPT: T5 Question and Answering

```
!nvidia-smi
```

Wed Dec 20 16:45:04 2023

NVIDIA-SMI 535.104.05				Driver Version: 535.104.05		CUDA Version	
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile U	
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	(
=====							
0	Tesla T4		Off	00000000:00:04.0	Off		
N/A	69C	P8	11W / 70W		0MiB / 15360MiB	0%	

Processes:							
GPU	GI	CI	PID	Type	Process name		(
	ID	ID					
=====							
No running processes found							

```

!pip install transformers==4.1.1
!pip install pytorch-lightning==1.1.1
!pip install tokenizers==0.9.4

```

```
!pip install sentencepiece==0.1.94
```

```
import torchtext
```

```
!pip install torchtext.data
```

```
ERROR: Could not find a version that satisfies the requirement torchtext.data
ERROR: No matching distribution found for torchtext.data
```

```
#import pytorch_lightning as pl
```

```
import re
```

```
import pandas as pd
```

```
import numpy as np
```

```
from transformers import (AdamW, T5ForConditionalGeneration, T5Tokenizer, get_lin
```

```
tokenizer = T5Tokenizer.from_pretrained("t5-base")
```

```
  spiece.model:   0%|          | 0.00/792k [00:00<?, ?B/s]
```

```
  tokenizer.json:                                1.39M/1.39M [00:00<00:00,
```

```
  100%                                                5.40MB/s]
```

```
  config.json: 100%                                1.21k/1.21k [00:00<00:00, 70.5kB/s]
```

```
/usr/local/lib/python3.10/dist-packages/transformers/models/t5/tokenization_t!
```

```
For now, this behavior is kept to avoid breaking backwards compatibility when
```

```
- Be aware that you SHOULD NOT rely on t5-base automatically truncating your :
```

```
- If you want to encode/pad to sequences longer than 512 you can either instai
```

```
- To avoid this warning, please instantiate this tokenizer with `model_max_le
```

```
  warnings.warn(
```

```
You are using the default legacy behaviour of the <class 'transformers.models
```

```
Special tokens have been added in the vocabulary, make sure the associated wo
```

```
sample = tokenizer("This is just trying this out", "another try at this")
```

```
sample_q = "trying this out again, will it work?"
```

```
sample.keys()
```

```
dict_keys(['input_ids', 'attention_mask'])
```

```
print(sample["input_ids"])
```

```
[100, 19, 131, 1119, 48, 91, 1, 430, 653, 44, 48, 1]
```

```
print(sample["attention_mask"])
```

```
[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]
```

```
pred = [tokenizer.decode(input, skip_special_tokens=True, clean_up_tokenization_s
```

```
" " join(pred)
```

```
.join(precu,
```

```
'This is just trying this out </s> another try at this </s>'
```

```
dict = {'questions': all_questions, 'answers': all_answers}
```

```
df = pd.DataFrame(dict)
```

```
print(df)
```

```

                                questions \
0      Just your thoughts on the court out there the ...
1      BI you just mentioned how you guys aren't on n...
2      You mentioned that you guys like to have fun. ...
3      CJ had mentioned the prize money a couple time...
4      What's the best and worst thing you can say ab...
...
1035   You had many ups and downs this season.  What'...
1036   George can you talk about the long touchdown r...
1037   George and Andrew I'm curious this last month ...
1038   George probably played his last game as a Bron...
1039   How do you want this team to be remembered?  H...

                                answers
0      My thoughts on the court is it looks like a...
1      Goofy.  We come to work we all have fun.  0...
2      No we ain't thought about it but I know it'...
3      For me I think my motivating factor is just...
4      I don't know you caught me off guard with t...
...
1035   Definitely my favorite part of the journey ...
1036   Yes in the first half we kicked a couple of...
1037   You know how strong the team is the perseve...
1038   George's legacy has been imprinted for a lo...
1039   I'm excited.  I believe God is up to someth...

```

```
[1040 rows x 2 columns]
```

```
import torch
```

```
from torch.utils.data import Dataset, DataLoader
```

```
from sklearn.model_selection import train_test_split
```

```
class t5ds(Dataset):
```

```
    def __init__(self, data, tokenizer: T5Tokenizer, q_len: int = 396, a_len: int =
```

```
        self.tokenizer = tokenizer
```

```
        self.data = data
```

```
        self.q_len = q_len
```

```
        self.a_len = a_len
```

```
    def __len__(self):
```

```
        return len(self.data)
```


[illegible]

[illegible]

```
from sklearn.model_selection import train_test_split
train, test = train_test_split(df, test_size=.1)
```

```
train.shape, test.shape
((936, 2), (104, 2))
```

```
!pip install pytorch-lightning==1.5.10
!pip install omegaconf -U
!pip install hydra-core --upgrade
import pytorch_lightning as pl
print(pl.__version__)
```

```
class t5DataModule(pl.LightningDataModule):
    def __init__(self, train_df, test_df, tokenizer, batch_size = 8, q_max: int = 3
        super().__init__()
        self.batch_size = batch_size
        self.test_df = test_df
        self.train_df = train_df
        self.tokenizer = tokenizer
        self.q_max = q_max
        self.a_max = a_max

    def setup(self):
        self.train_ds = t5ds(self.train_df, self.tokenizer, self.q_max, self.a_max)
        self.test_ds = t5ds(self.test_df, self.tokenizer, self.q_max, self.a_max)

    def train_dataloader(self):
        return DataLoader(self.train_ds, batch_size = self.batch_size, shuffle=True,

    def val_dataloader(self):
        return DataLoader(self.train_ds, batch_size = 1, num_workers=4)

    def test_dataloader(self):
        return DataLoader(self.test_ds, batch_size = 1, num_workers=4)
```

```
BATCH_SIZE = 2
N_EPOCHS = 4
```

```
data_module = t5DataModule(train, test, tokenizer, batch_size = BATCH_SIZE)
data_module.setup()
```

```
model = T5ForConditionalGeneration.from_pretrained('t5-base', return_dict=True)
```

```
model.config
```

```
T5Config {
  "_name_or_path": "t5-base",
  "architectures": [
    "T5ForConditionalGeneration"
  ],
  "classifier_dropout": 0.0,
  "d_ff": 3072,
  "d_kv": 64,
  "d_model": 768,
  "decoder_start_token_id": 0,
  "dense_act_fn": "relu",
  "dropout_rate": 0.1,
  "eos_token_id": 1,
  "feed_forward_proj": "relu",
  "initializer_factor": 1.0,
  "is_encoder_decoder": true,
  "is_gated_act": false,
  "layer_norm_epsilon": 1e-06,
  "model_type": "t5",
  "n_positions": 512,
  "num_decoder_layers": 12,
  "num_heads": 12,
  "num_layers": 12,
  "output_past": true,
  "pad_token_id": 0,
  "relative_attention_max_distance": 128,
  "relative_attention_num_buckets": 32,
  "task_specific_params": {
    "summarization": {
      "early_stopping": true,
      "length_penalty": 2.0,
      "max_length": 200,
      "min_length": 30,
      "no_repeat_ngram_size": 3,
      "num_beams": 4,
      "prefix": "summarize: "
    },
    "translation_en_to_de": {
      "early_stopping": true,
      "max_length": 300,
      "num_beams": 4,
      "prefix": "translate English to German: "
```

```

    },
    "translation_en_to_fr": {
        "early_stopping": true,
        "max_length": 300,
        "num_beams": 4,
        "prefix": "translate English to French: "
    },
    "translation_en_to_ro": {
        "early_stopping": true,
        "max_length": 300,
        "num_beams": 4,
        "prefix": "translate English to Romanian: "
    }
},
"transformers_version": "4.35.2",
"use_cache": true,

```

```

class newT5(pl.LightningModule):
    def __init__(self):
        super().__init__()
        self.model = T5ForConditionalGeneration.from_pretrained('t5-base', return_dict=True)

    def forward(self, input_ids, attention_mask, labels=None):
        output = self.model(input_ids=input_ids, attention_mask=attention_mask, labels=labels)
        if labels is not None:
            return output.loss, output.logits
        else:
            return output.logits
        return output.loss, output.logits

    def training_step(self, batch, batch_idx):
        input_ids = batch["input_ids"]
        attention_mask = batch["attention_mask"]
        labels = batch["labels"]
        loss, outputs = self(input_ids, attention_mask, labels)
        self.log("train_loss", loss, prog_bar=True, logger=True)
        return loss

    def validation_step(self, batch, batch_idx):
        input_ids = batch["input_ids"]
        attention_mask = batch["attention_mask"]
        labels = batch["labels"]
        loss, outputs = self(input_ids, attention_mask, labels)
        self.log("val_loss", loss, prog_bar=True, logger=True)
        return loss

    def test_step(self, batch, batch_idx):
        input_ids = batch["input_ids"]
        attention_mask = batch["attention_mask"]
        labels = batch["labels"]
        loss, outputs = self(input_ids, attention_mask, labels)
        self.log("test_loss", loss, prog_bar=True, logger=True)

```

```

self.log('test_loss', loss, prog_bar=True, logger=True)
return loss

def configure_optimizers(self):
    return AdamW(self.parameters(), lr=.0001)

model = newT5()

from keras.callbacks import ModelCheckpoint
checkpoint_callback = ModelCheckpoint(filepath="checkpoints", filename="best_checkpoint.h5")

trainer = pl.Trainer(checkpoint_callback=checkpoint_callback, max_epochs = 3, gpus=-1,
                      callbacks=[ModelCheckpoint(filepath="checkpoints", filename="best_checkpoint.h5")],
                      reload_modules=[
                          /usr/local/lib/python3.10/dist-packages/pytorch_lightning/trainer/connectors/connector.py,
                          rank_zero_deprecation(
                          /usr/local/lib/python3.10/dist-packages/pytorch_lightning/trainer/connectors/connector.py,
                          rank_zero_deprecation(
                          INFO:pytorch_lightning.utilities.distributed:GPU available: True, used: True
                          INFO:pytorch_lightning.utilities.distributed:TPU available: False, using: 0 TPU
                          INFO:pytorch_lightning.utilities.distributed:IPU available: False, using: 0 IPU

%load_ext tensorboard

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

%tensorboard --logdir ./lightning_logs

```

TensorBoard**INACTIVE**

No dashboards are active for the current data set.

Probable causes:

- You haven't written any data to your event files.
- TensorBoard can't find your event files.

If you're new to using TensorBoard, and want to find out how to add data and set up your event files, check out the [README](#) and perhaps the [TensorBoard tutorial](#).

If you think TensorBoard is configured properly, please see [the section of the README devoted to missing data problems](#) and consider filing an issue

[Note: This code is for testing data problems and generating an issue](#)
on GitHub.

Last reload: Dec 23, 2023, 10:33:25 PM

Log directory: ./lightning_logs

```

trainer.fit(model, data_module)
/usr/local/lib/python3.10/dist-packages/pytorch_lightning/core/datamodule.py:
rank_zero_deprecation(
INFO:pytorch_lightning.accelerators.gpu:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES:
/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411: Futu
warnings.warn(
INFO:pytorch_lightning.callbacks.model_summary:
  | Name   | Type                               | Params
-----
0 | model  | T5ForConditionalGeneration        | 222 M
-----
222 M      Trainable params
0          Non-trainable params
222 M      Total params
891.614    Total estimated model params size (MB)
Validation sanity check: 0%                                0/2 [00:00<?, ?it/s]
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: U
warnings.warn(_create_warning_msg(
-----
RecursionError                                Traceback (most recent call last)
<ipython-input-61-7b6b8391c42e> in <cell line: 1>()
----> 1 trainer.fit(model, data_module)
-----
20 frames
... last 3 frames repeated, from the frame below ...

```

```

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py in
_wrapped_call_impl(self, *args, **kwargs)
    1516         return self._compiled_call_impl(*args, **kwargs) # type:
ignore[misc]
    1517         else:
-> 1518         return self._call_impl(*args, **kwargs)
    1519
    1520     def _call_impl(self, *args, **kwargs):

```

RecursionError: maximum recursion depth exceeded while calling a Python object

SEARCH STACK OVERFLOW

✓ RNN for Q and A

```

import nltk
nltk.download('punkt')
from nltk.tokenize import word_tokenize

token_q = []
token_a = []

for sent in all_questions:
    token_q.append(word_tokenize(sent))

for sent in all_answers:
    token_a.append(word_tokenize(sent))

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!

print(token_q[0])
print(token_a[0])

['Just', 'your', 'thoughts', 'on', 'the', 'court', 'out', 'there', 'the', 'vei
['My', 'thoughts', 'on', 'the', 'court', 'is', 'it', 'looks', 'like', 'a', 's'

def maxi(sents, others):
    for x in range(100):
        nmax = 0
        to_remove = None
        for sent in sents:
            if len(sent) > nmax:
                nmax = len(sent)
                to_remove = sents.index(sent)
        del sents[to_remove]
        del others[to_remove]
    return others, sents

```

```
def bracket_questions(sent):
    sent = ['<q>'] + sent + ['</q>']
    return sent

def bracket_answer(sent):
    sent = ['<a>'] + sent + ['</a>']
    return sent

print(len(token_q))
print(len(token_a))
token_q, token_a = maxi(token_a, token_q)
print(len(token_q))
print(len(token_a))

for x in range(len(token_q)):
    token_q[x] = bracket_questions(token_q[x])
    token_a[x] = bracket_answer(token_a[x])

    1040
    1040
    940
    940

from google.colab import drive
drive.mount('/content/drive')

data_dir = 'drive/MyDrive/'
    Mounted at /content/drive

print(token_q[0])
print(token_a[0])

    ['<q>', 'Just', 'your', 'thoughts', 'on', 'the', 'court', 'out', 'there', 'the']
    ['<a>', 'My', 'thoughts', 'on', 'the', 'court', 'is', 'it', 'looks', 'like',

vocab = {}
i = 0
for sent in token_q:
    for token in sent:
        if token not in vocab:
            vocab[token] = i
            i += 1

for sent in token_a:
    for token in sent:
        if token not in vocab:
            vocab[token] = i
            i += 1
```



```

idx_to_token = {idx: token for token, idx in vocab.items()}

print(vocab)

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import numpy as np

def load_glove_embeddings(glove_file):
    weights_matrix = torch.zeros(len(vocab), 100)
    embeddings = {}
    with open(glove_file, 'r', encoding='utf-8') as file:
        for line in file:
            values = line.split()
            word = values[0]
            vector = np.array([float(val) for val in values[1:]])
            embeddings[word] = vector
            if word in vocab:
                idx = vocab[word]
                weights_matrix[idx] = torch.tensor(np.asarray(values[1:], "float32"))
    for word, idx in vocab.items():
        if not torch.any(weights_matrix[idx]):
            weights_matrix[idx] = torch.normal(0.0, 0.6, size=(100, ))

    return weights_matrix, embeddings

weights_matrix, glove_embeddings = load_glove_embeddings(data_dir + './glove.6B.1

all_tokens = []
for i in range(len(token_q)):
    for token in token_q[i]:
        if token == 'J':
            print(all_questions[i])
            all_tokens.append(token)
    for token in token_a[i]:
        if token == 'J':
            print(all_answers[i])
            all_tokens.append(token)

encode = np.vectorize(lambda w: vocab[w])
encoded = encode(all_tokens)

sequences = []
goals = []

for i in range(len(encoded) - 41):
    sequences.append((encoded[i:i+40]))

```

```
goals.append((encoded[1+1:1+41]))
print(sequences[0])
print(goals[0])
```

[0	1	2	3	4	5	6	7	8	5	9	10	2	3
	4	11	12	13	5	14	15	13	16	17	18	19	20	3186
	3187	3	4	5	6	88	91	1557	62	31	213	19]		
[1	2	3	4	5	6	7	8	5	9	10	2	3	4
	11	12	13	5	14	15	13	16	17	18	19	20	3186	3187
	3	4	5	6	88	91	1557	62	31	213	19	1586]		

```
class NEW_DS(Dataset):
```

```
    def __init__(self, X, Y):
        self.X = X # Convert X to a PyTorch LongTensor
        self.Y = Y
```

```
    def __len__(self):
        return len(self.X)
```

```
    def __getitem__(self, idx):

        return torch.tensor(self.X[idx]), torch.tensor(self.Y[idx])
```

```
batch_size = 32
dataset = NEW_DS(sequences, goals)
```

```
data_loader = DataLoader(dataset, batch_size = 32, shuffle=False)
```

```
class MyLSTM(nn.Module):
    def __init__(self, vocab_size, embedding_size, hidden_size, embedding_weights):
        super(MyLSTM, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_size)
        self.embedding.weight = nn.Parameter(embedding_weights)
        self.embedding.weight.requires_grad = True
        self.lstm = nn.LSTM(embedding_size, hidden_size, batch_first=True)
        self.linear = nn.Linear(hidden_size, vocab_size)

    def forward(self, x):
        embedded = self.embedding(x)
        lstm_out, _ = self.lstm(embedded)
        output = self.linear(lstm_out)
        return output
```

```
import matplotlib.pyplot as plt
```

```
vocab_size = len(vocab)
embedding_size = 100
hidden_size = 256
```

```

model = MyLSTM(vocab_size, embedding_size, hidden_size, weights_matrix)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# Training loop
num_epochs = 4
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)
y = []
for epoch in range(num_epochs):
    model.train()
    total_loss = 0

    for batch_idx, (inputs, targets) in enumerate(data_loader):
        inputs, targets = inputs.to(device), targets.to(device)

        optimizer.zero_grad()
        outputs = model(inputs)

        loss = criterion(outputs.permute(0, 2, 1), targets) # Permute outputs to
        total_loss += loss.item()

        loss.backward()
        optimizer.step()

        if batch_idx % 100 == 0:
            print(f'Epoch [{epoch+1}/{num_epochs}], Batch [{batch_idx+1}/{len(data_loader)}], Loss: {loss.item():.4f}')

    average_loss = total_loss / len(data_loader)
    print(f'Epoch [{epoch+1}/{num_epochs}], Average Loss: {average_loss:.4f}')
    y.append(average_loss)

plt.xlabel('Epochs')
plt.ylabel('Avg Loss Per Epoch')
plt.title('Loss')

x = list(range(1, num_epochs + 1))

plt.plot(x, y)

plt.show()

```

Epoch [1/4], Batch [1/2553], Loss: 8.4941
Epoch [1/4], Batch [101/2553], Loss: 6.0208
Epoch [1/4], Batch [201/2553], Loss: 6.0986
Epoch [1/4], Batch [301/2553], Loss: 6.0233
Epoch [1/4], Batch [401/2553], Loss: 5.3439
Epoch [1/4], Batch [501/2553], Loss: 4.9774
Epoch [1/4], Batch [601/2553], Loss: 5.0792
Epoch [1/4], Batch [701/2553], Loss: 5.1037
Epoch [1/4], Batch [801/2553], Loss: 5.2506
Epoch [1/4], Batch [901/2553], Loss: 5.0924

```
Epoch [1/4], Batch [1001/2553], Loss: 4.9033
Epoch [1/4], Batch [1101/2553], Loss: 5.4155
Epoch [1/4], Batch [1201/2553], Loss: 4.6104
Epoch [1/4], Batch [1301/2553], Loss: 4.4934
Epoch [1/4], Batch [1401/2553], Loss: 4.8183
Epoch [1/4], Batch [1501/2553], Loss: 4.4867
Epoch [1/4], Batch [1601/2553], Loss: 4.4682
Epoch [1/4], Batch [1701/2553], Loss: 4.3999
Epoch [1/4], Batch [1801/2553], Loss: 4.4497
Epoch [1/4], Batch [1901/2553], Loss: 5.8747
Epoch [1/4], Batch [2001/2553], Loss: 5.1210
Epoch [1/4], Batch [2101/2553], Loss: 4.6843
Epoch [1/4], Batch [2201/2553], Loss: 6.1195
Epoch [1/4], Batch [2301/2553], Loss: 4.0410
Epoch [1/4], Batch [2401/2553], Loss: 4.5573
Epoch [1/4], Batch [2501/2553], Loss: 4.6640
Epoch [1/4], Average Loss: 5.1581
Epoch [2/4], Batch [1/2553], Loss: 4.5541
Epoch [2/4], Batch [101/2553], Loss: 4.3439
Epoch [2/4], Batch [201/2553], Loss: 4.3460
Epoch [2/4], Batch [301/2553], Loss: 4.5638
Epoch [2/4], Batch [401/2553], Loss: 3.8549
Epoch [2/4], Batch [501/2553], Loss: 4.0581
Epoch [2/4], Batch [601/2553], Loss: 4.0330
Epoch [2/4], Batch [701/2553], Loss: 4.0717
Epoch [2/4], Batch [801/2553], Loss: 4.5112
Epoch [2/4], Batch [901/2553], Loss: 4.3105
Epoch [2/4], Batch [1001/2553], Loss: 4.2491
Epoch [2/4], Batch [1101/2553], Loss: 4.5146
Epoch [2/4], Batch [1201/2553], Loss: 4.0246
Epoch [2/4], Batch [1301/2553], Loss: 3.6104
Epoch [2/4], Batch [1401/2553], Loss: 3.9061
Epoch [2/4], Batch [1501/2553], Loss: 3.7914
Epoch [2/4], Batch [1601/2553], Loss: 3.9098
Epoch [2/4], Batch [1701/2553], Loss: 3.6242
Epoch [2/4], Batch [1801/2553], Loss: 3.9324
Epoch [2/4], Batch [1901/2553], Loss: 4.6476
Epoch [2/4], Batch [2001/2553], Loss: 4.6767
Epoch [2/4], Batch [2101/2553], Loss: 4.2307
Epoch [2/4], Batch [2201/2553], Loss: 5.2817
Epoch [2/4], Batch [2301/2553], Loss: 3.6146
Epoch [2/4], Batch [2401/2553], Loss: 4.0233
Epoch [2/4], Batch [2501/2553], Loss: 4.0975
Epoch [2/4], Average Loss: 4.2554
Epoch [3/4], Batch [1/2553], Loss: 3.9095
Epoch [3/4], Batch [101/2553], Loss: 4.0255
Epoch [3/4], Batch [201/2553], Loss: 3.8844
Epoch [3/4], Batch [301/2553], Loss: 4.0980
```

```
model.eval()
```

```
# Define a starting sequence for prediction
```

```
start_sequence = ['<q>', 'how', 'was', 'the', 'game', '</q>', '<a>']
```

```
.. - . . . . .
```

```
# Encode the starting sequence to integers
encoded_start_sequence = [vocab[char] for char in start_sequence]

# Convert the encoded start sequence to a PyTorch tensor
inputs = torch.tensor(encoded_start_sequence).unsqueeze(0) # Add batch dimension

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
inputs = inputs.to(device)

# Predict the next character
with torch.no_grad():
    hidden = None
    for i in range(20): # Predict the next 50 characters
        #print(inputs.shape)
        outputs = model(inputs)
        # Get the output logits for the last character in the sequence
        last_output = outputs[:, -1, :]
        # Get the index of the predicted character
        predicted_index = torch.argmax(last_output, dim=1).item()
        # Map the index back to the character
        predicted = idx_to_token[predicted_index]
        # Print the predicted character
        if predicted == '</a>':
            break
        print(" " + predicted + " ", end='')

    # Update the input sequence for the next prediction
    inputs = torch.cat((inputs, torch.tensor([[predicted_index]]).to(device)))

Yeah .
```

